



CUDA graph replay: $5.13\times$ faster · $0.480 \rightarrow 2.463$ recovered B/s

● Reference ● CUDA graph ● Repetition 1 ■ Repetition 2 ▲ Repetition 3


A Encode · three repetitions each



B Decode · three repetitions each



C Useful throughput · encode + decode



D GPU memory · distinct measures

