

Carrier / method	Frozen score	AUC [grouped 95% interval]	Scorable stego	Unique matched controls	Failed stego included
UTF-8 / Fixed	Mean surprisal	1.0000 [1.0000, 1.0000]	20	20	0
UTF-8 / Fixed	Mean log2 rank	1.0000 [1.0000, 1.0000]	20	20	0
UTF-8 / Gated	Mean surprisal	1.0000 [1.0000, 1.0000]	20	20	0
UTF-8 / Gated	Mean log2 rank	1.0000 [1.0000, 1.0000]	20	20	0
UTF-8 / A1	Mean surprisal	0.3800 [0.2600, 0.5025]	20	20	20
UTF-8 / A1	Mean log2 rank	0.3600 [0.2425, 0.4801]	20	20	20
PNG / Fixed	Mean surprisal	0.8850 [0.7425, 0.9850]	20	20	0
PNG / Fixed	Mean log2 rank	0.8400 [0.6825, 0.9650]	20	20	0
PNG / Gated	Mean surprisal	0.8300 [0.7000, 0.9375]	20	20	0
PNG / Gated	Mean log2 rank	0.7700 [0.6299, 0.8975]	20	20	0
PNG / A1	Mean surprisal	0.5575 [0.3750, 0.7426]	20	20	0
PNG / A1	Mean log2 rank	0.5750 [0.4050, 0.7576]	20	20	0

Exploratory known-model/context detectability. Both frozen whole-carrier scores use higher-is-stego direction, including AUCs below 0.5. Intervals reuse the accepted 2,000-draw payload-group bootstrap with class/length strata and shared controls. There are 40 independent ordinary traces overall, not 120 per-method control replicates or 240 per-score observations. Each cell has 20 unique matched artifacts; no V2 duplicates or unscorable carriers were dropped. All 20 failed A1 text carriers are included. Small empirical perfect separation and a [1,1] bootstrap interval do not prove perfect population detection; near-chance scores do not prove security.